

Data Analysis using R

Importing

Sven Werenbeck-Ueding

21.09.2023

Source: [Wickham and Grolemund \(2016\)](#)

Popular file formats

Excel Spreadsheet (.xlsx)

- Format for storing human-readable, rectangular data
- One or more sheets containing rectangular data
- Older spreadsheet file format: .xls

Popular file formats

Comma Separated Values (.csv)

- Format for storing human-readable, rectangular data
- Values are most commonly separated by , or ;
- Each line corresponds to a row
- First line usually contains column names

Example Data

- Mexican survey data used in [Feigenberg \(2020a\)](#) and provided by [Feigenberg \(2020b\)](#)
 - Encuesta Nacional de Ocupación y Empleo (ENOE)
 - Q3 2003 to Q3 2013
 - Quarterly rotating panel: Households included for 5 quarters
 - Records whether any household member leaves to the US
 - Potential migrants are restricted from ages 15 to 65
 - Explanatory variables: age, gender, marital status and education for all household members
- In folder `data/raw/` of the online repository in csv and excel format
- Data was altered to satisfy needs for this course

Importing Data

readxl



The readxl package makes it easy to get data out of Excel and into R. [...] readxl has no external dependencies, so it's easy to install and use on all operating systems. It is designed to work with tabular data. readxl supports both the legacy .xls format and the modern xml-based .xlsx format.

Wickham and Bryan (2022)



readxl is part of the tidyverse but not part of its core packages!

Load the package to use it:

```
#install.packages("readxl")  
library(readxl)
```


read_excel()

```
?read_excel
```

The `read_excel()` function is capable of reading data from `.xls` and `.xlsx` files. If the file format is known, it is recommended to directly use `read_xls()` or `read_xlsx()`, respectively, to prevent R from guessing.

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

path

The path to the file you want to read

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

sheet

String or integer specifying the sheet name or position to read; defaults to the first sheet

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

range

Excel expression for the cell range to read from, e. g. "B3:D87". Can also be used to specify the sheet name like "Budget!B2:G14".

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

col_names

- Either TRUE/FALSE or a character vector of column names
- If TRUE, the default, the first row is used as column names

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

col_types

read_excel()

```
?read_excel
```

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

na

Character vector to interpret as missing values; defaults to empty cells, i. e. ""

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

skip

Number of rows to skip before reading the data; defaults to 0

read_excel()

?read_excel

Function call and arguments:

```
read_excel(  
  path,  
  sheet = NULL,  
  range = NULL,  
  col_names = TRUE,  
  col_types = NULL,  
  na = "",  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

n_max

Sets the maximum number of rows to read; defaults to Inf

Let's try this out!

Code

- Import raw data on migration choice of Mexicans and their socioeconomic characteristics
- Data sets can be found in the data folder of the repository

```
read_excel("data/raw/enoe/enoe.xlsx")
```

Message

```
## New names:
## • `` -> `...2`
## • `` -> `...3`
## • `` -> `...4`
## • `` -> `...5`
## • `` -> `...6`
## • `` -> `...7`
## • `` -> `...8`
## • `` -> `...9`
## • `` -> `...10`
## • `` -> `...11`
## • `` -> `...12`
```



What happened?

Let's take a look at the excel file:

What happened?

Let's take a look at the excel file:

Another try...

Code

```
read_excel("data/raw/enoe/enoe.xlsx",  
            skip = 3, # Skip the first three lines  
            na = c("", "N/A")) # Interpret empty cells and cells containing "N/A" as missing
```

Message

Output

```
## # A tibble: 165,457 × 12  
##   id      migrate age  munici...1 fence year quarter sex  marit...2 empl_...3 educ  
##   <chr> <chr>   <chr> <chr>      <chr> <chr> <chr> <chr> <chr>   <chr>   <chr>  
## 1 189889 No      50    2004      0    2004 Q3     Fema... Single Unempl... 12  
## 2 189889 No      50    2004      0    2004 Q4     Fema... Single Unempl... 12  
## 3 189889 No      50    2004      0    2005 Q1     Fema... Single Unempl... 12  
## 4 189890 No      26    2004      0    2005 Q4     Male   Married Full-t... 10  
## 5 189890 No      26    2004      0    2006 Q1     Male   Married Full-t... 10  
## 6 189890 No      26    2004      0    2006 Q2     Male   <NA>    Full-t... 10  
## 7 189891 No      36    2004      0    2006 Q4     Male   Married Full-t... 6  
## 8 189891 No      36    2004      0    2007 Q1     Male   Married Full-t... 6  
## 9 189891 No      36    2004      0    2007 Q2     Male   Married Full-t... 6
```

readr



The goal of `readr` is to provide a fast and friendly way to read rectangular data from delimited files, such as comma-separated values (CSV) and tab-separated values (TSV). It is designed to parse many types of data found in the wild, while providing an informative problem report when parsing leads to unexpected results.

Wickham, Hester, and Bryan (2022)

read_csv()

```
?read_csv
```

read_csv() function is a special case of `read_delim()` with the **separator set to ,** used for reading files with comma separated values. The related **read_csv2()**, in contrast, **assumes a ; as separator** for values and interprets `,` as decimal point (which is common in some European countries).

- Considerably faster than the base R solution `read.csv()`
- Consistent parameter naming (`col_names` and `col_type` instead of `header` and `colClasses`)
- Automatically parses common date formats, but leaves strings unaltered
- Progression bar for big data sets

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

file

The path to the file you want to read

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

col_names

- Either TRUE/FALSE or a character vector of column names
- If TRUE, the default, the first row is used as column names
- If FALSE, column names are generated in this fashion: X1, X2, \dots

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

col_types

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

col_select

Character vector containing columns to include in the resulting data set; defaults to NULL

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

na

Character vector to interpret as missing values; defaults to `c("", "NA")`

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

skip

Number of lines to skip before reading the data; defaults to 0

read_csv()

?read_csv

Function call and arguments:

```
read_csv(  
  file,  
  col_names = TRUE,  
  col_types = NULL,  
  col_select = NULL,  
  na = c("", "NA"),  
  skip = 0,  
  n_max = Inf,  
  ...  
)
```

n_max

Sets the maximum number of lines to read; defaults to `Inf`

Import the ENE/ENOE Data Set from CSV

Code

```
read_csv("data/raw/enoe/enoe.csv")
```

Message

```
## New names:
## • `` -> `...2`
## • `` -> `...3`
## • `` -> `...4`
## • `` -> `...5`
## • `` -> `...6`
## • `` -> `...7`
## • `` -> `...8`
## • `` -> `...9`
## • `` -> `...10`
## • `` -> `...11`
## • `` -> `...12`
```

```
## Warning: One or more parsing issues, call `problems()` on your data frame for details,
## e.g.:
```




Similar problems as before!

- Skip first three lines
- Set first line (after skipping three lines) as column names
- Interpret N/A as missing

Import ENE/ENOE Data with Function Arguments

Code

```
read_csv("data/raw/enoe/enoe.csv",  
         skip = 3,  
         col_names = TRUE, # The default  
         na = c("", "N/A"))
```

Message

```
## Rows: 165457 Columns: 12  
## — Column specification —————  
## Delimiter: ","  
## chr (11): migrate, age, municipality, fence, year, quarter, sex, marital_sta...  
## dbl (1): id  
##  
## i Use `spec()` to retrieve the full column specification for this data.  
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.
```

Output

```
## # A tibble: 165,457 × 12
```

Data Types

Vectors in R

- Vectors can be distinguished in atomic vectors and lists
 - Atomic vector: All elements must have the same type
 - Lists: Elements can have different types
- If you are already familiar with R, you probably have encountered vectors on several occasions, e. g. the sequence `1:10` is an atomic vector containing all integers between 1 and 10

```
1:10
```

```
## [1] 1 2 3 4 5 6 7 8 9 10
```

```
typeof(1:10) # Check the type of the integer vector 1:10
```

```
## [1] "integer"
```

Primary Types of Atomic Vectors

Type	Description	How to check
Logical	TRUE/FALSE or their abbreviations T/F. Simple calculations can be conducted on vectors of this type, e.g. <code>sum(T, T, F, T)</code> gives 3 because TRUE is interpreted as 1 and FALSE as 0.	<code>is.logical()</code>
Character	Strings surrounded by " or '	<code>is.character()</code>
Double	Numerical values with decimals. Special values are <code>Inf</code> , <code>-Inf</code> and <code>NaN</code> .	<code>is.double()</code>
Integer	Numerical values that cannot contain fractional values. Must be followed by <code>L</code>	<code>is.integer()</code>



Both, Integers and doubles, are numerical values!

`is.numeric()` returns TRUE for values that are either integers *or* doubles.

Factors: Representation of Categorical Data

- Factors contain predefined values only
- Integer vector with attributes `class` ("factor") and `levels` (set of values)

```
# Create a factor vector with two levels c("Employed", "Unemployed")
fct <- factor(c("Employed", "Employed", "Unemployed", "Employed"))

fct
```

```
## [1] Employed   Employed   Unemployed Employed
## Levels: Employed Unemployed
```

```
# Take a look at the vector
str(fct)
```

```
## Factor w/ 2 levels "Employed","Unemployed": 1 1 2 1
```

Factors: Representation of Categorical Data

```
# Check the vector type  
typeof(fct)
```

```
## [1] "integer"
```

```
# Get the vector attributes  
attributes(fct)
```

```
## $levels  
## [1] "Employed" "Unemployed"  
##  
## $class  
## [1] "factor"
```

```
# To get the levels of a vector, use the short-hand levels() function instead  
levels(fct)
```

```
## [1] "Employed" "Unemployed"
```

Lists

```
# Construct a list
our_list <- list(1:10, # Integer sequence
               seq(0, 1, by = 0.25), # Double sequence
               c("This", "is", "a", "character", "vector"), # Character vector
               factor(c("Employed", "Employed", "Unemployed", "Employed"))) # Factor vector
```

```
typeof(our_list)
```

```
## [1] "list"
```

```
# Take a look at the list structure
str(our_list)
```

```
## List of 4
## $ : int [1:10] 1 2 3 4 5 6 7 8 9 10
## $ : num [1:5] 0 0.25 0.5 0.75 1
## $ : chr [1:5] "This" "is" "a" "character" ...
## $ : Factor w/ 2 levels "Employed","Unemployed": 1 1 2 1
```


Lists



Lists may contain more complex objects than atomic vectors, such as lists.

```
list(1:10,  
      seq(0, 1, by = 0.25),  
      list("a", 1:2))
```

```
## [[1]]  
## [1]  1  2  3  4  5  6  7  8  9 10  
##  
## [[2]]  
## [1] 0.00 0.25 0.50 0.75 1.00  
##
```

Data Frames

- Data sets are usually represented as `data.frame` objects in R
- A `data.frame` is essentially a named list of vectors with equal length

```
# Construct a data frame
df <- data.frame(income = c(0, 500, 3000),
                 empl_status = factor(c("Unemployed", "Employed", "Employed")))
```

Data Frames

```
# Check the type  
typeof(df)
```

```
## [1] "list"
```

```
# Get the attributes  
attributes(df)
```

```
## $names  
## [1] "income"      "empl_status"  
##  
## $class  
## [1] "data.frame"  
##  
## $row.names  
## [1] 1 2 3
```

tibble



A `tibble`, or `tbl_df`, is a modern reimagining of the `data.frame`, keeping what time has proven to be effective, and throwing out what is not. Tibbles are `data.frames` that are lazy and surly: they do less (i.e. they don't change variable names or types, and don't do partial matching) and complain more (e.g. when a variable does not exist). This forces you to confront problems earlier, typically leading to cleaner, more expressive code. Tibbles also have an enhanced `print()` method which makes them easier to use with large datasets containing complex objects.

Müller and Wickham (2022)

Data Frames vs. Tibbles

Tibbles' enhanced `print()` function shows only the first 10 rows and displays information on the data structure:

```
# Construct a data frame
df <- data.frame(
  income = c(0, 500, 3000),
  empl_status = factor(
    c("Unemployed", "Employed", "Employed")
  )
)

# Print the data frame
df
```

```
##   income empl_status
## 1      0  Unemployed
## 2    500    Employed
## 3   3000    Employed
```

```
# Construct a tibble
tbl <- tibble(
  income = c(0, 500, 3000),
```

Data Frames vs. Tibbles

When subsetting, `tibble` gives a warning if the column does not exist:

```
df$gender
```

```
## NULL
```

```
tbl$gender
```

```
## Warning: Unknown or uninitialised column: `gender`.
```

```
## NULL
```

Importing Data as Tibble

`read_excel()` and `read_csv()` automatically create a tibble object:

```
data <- read_csv("data/raw/enoe/enoe.csv", skip = 3, col_names = TRUE, na = c("", "N/A"))  
typeof(data)
```

```
## Rows: 165457 Columns: 12  
## — Column specification —————  
## Delimiter: ","  
## chr (11): migrate, age, municipality, fence, year, quarter, sex, marital_sta...  
## dbl (1): id  
##  
## i Use `spec()` to retrieve the full column specification for this data.  
## i Specify the column types or set `show_col_types = FALSE` to quiet this message.  
  
## [1] "spec_tbl_df" "tbl_df"      "tbl"         "data.frame"
```



Further Literature

- Further information on tibbles can be found in chapter 10 of [Wickham and Grolemund \(2016\)](#) and the package documentation (see [Müller and Wickham \(2022\)](#)).
- Chapter 3 of [Wickham \(2019\)](#) discusses the technical properties of vectors in R more deeply.

Exporting Data

write_csv()

```
?write_csv
```

write_csv() function is a special case of `write_delim()` with the **separator set to ,**. Again, there is a related **write_csv2()** which sets the separator to `;`.

- Considerably faster than the base R solution `write.csv()`
- Does not include row names as a column in the written file
- Progression bar for big data sets

write_csv()

?write_csv

Function call and arguments:

```
write_csv(  
  x,  
  file,  
  na = "NA",  
  append = FALSE,  
  col_names = !append,  
  ...  
)
```

x

A data.frame or tibble to write to .csv

write_csv()

?write_csv

Function call and arguments:

```
write_csv(  
  x,  
  file,  
  na = "NA",  
  append = FALSE,  
  col_names = !append,  
  ...  
)
```

file

File (path) to write to

write_csv()

?write_csv

Function call and arguments:

```
write_csv(  
  x,  
  file,  
  na = "NA",  
  append = FALSE,  
  col_names = !append,  
  ...  
)
```

na

String used for missing values; defaults to "NA"

write_csv()

?write_csv

Function call and arguments:

```
write_csv(  
  x,  
  file,  
  na = "NA",  
  append = FALSE,  
  col_names = !append,  
  ...  
)
```

append

If FALSE, the default, the existing will be overwritten. If TRUE, it will be appended to the existing file.

write_csv()

?write_csv

Function call and arguments:

```
write_csv(  
  x,  
  file,  
  na = "NA",  
  append = FALSE,  
  col_names = !append,  
  ...  
)
```

col_names

- If FALSE, column names will not be included at the top of the file.
- If TRUE, column names will be included.
- The default is to take the opposite value given to argument append

Export the ENE/ENOE Data Set to CSV

Code

```
data <- read_csv("data/raw/enoe/enoe.csv", skip = 3, col_names = TRUE, na = c("", "N/A"))  
write_csv(data, "data/raw/enoe.csv")
```

Output

The exported data is stored in csv format in the raw folder of the data folder:

```
list.files("data/raw/")
```

```
## [1] "eno"           "eno.csv"       "fence_construction"
```


References

Feigenberg, B. (2020a). "Fenced Out: The Impact of Border Construction on US-Mexico Migration". In: *American Economic Journal: Applied Economics* 12.3, pp. 106-39. DOI: [10.1257/app.20170231](https://doi.org/10.1257/app.20170231). URL: <https://www.aeaweb.org/articles?id=10.1257/app.20170231>.

Feigenberg, B. (2020b). *Replication package for: Fenced Out: The Impact of Border Construction on US-Mexico Migration*. <https://www.aeaweb.org/journals/dataset?id=10.1257/app.20170231>, Last accessed on 2022-12-29.

Müller, K. and H. Wickham (2022). *tibble: Simple Data Frames*. <https://tibble.tidyverse.org/>, <https://github.com/tidyverse/tibble>.

Wickham, H. (2019). *Advanced R*. 2nd. Chapman & Hall/CRC. URL: <http://adv-r.had.co.nz/>.

Wickham, H. and J. Bryan (2022). *readxl: Read Excel Files*. <https://readxl.tidyverse.org>, <https://github.com/tidyverse/readxl>.

Wickham, H. and G. Grolemund (2016). *R for data science. import, tidy, transform, visualize, and model data*. O'Reilly. URL: <https://r4ds.had.co.nz/>.

Wickham, H., J. Hester, and J. Bryan (2022). *readr: Read Rectangular Text Data*. <https://readr.tidyverse.org>, <https://github.com/tidyverse/readr>.