

# **Data Analysis using R**

## **Data Wrangling**

**Sven Werenbeck-Ueding**

**21.09.2023**

*Source:* [Wickham and Grolemund \(2016\)](#)

# Prerequisites

## Code

```
# Load the ENE/ENOE data set
tbl_enoe <- read_csv("data/raw/enoe/enoe.csv", skip = 3, na = c("", "N/A"))

tbl_enoe
```

## Output

```
## # A tibble: 165,457 × 12
##       id migrate age  municipi...1 fence year quarter sex  marit...2 empl_...3 educ
##   <dbl> <chr>   <chr> <chr>      <chr> <chr> <chr>  <chr> <chr>   <chr>   <chr>
## 1 189889 No      50    2004      0    2004 Q3     Fema... Single  Unempl... 12
## 2 189889 No      50    2004      0    2004 Q4     Fema... Single  Unempl... 12
## 3 189889 No      50    2004      0    2005 Q1     Fema... Single  Unempl... 12
## 4 189890 No      26    2004      0    2005 Q4     Male   Married Full-t... 10
## 5 189890 No      26    2004      0    2006 Q1     Male   Married Full-t... 10
## 6 189890 No      26    2004      0    2006 Q2     Male   <NA>    Full-t... 10
## 7 189891 No      36    2004      0    2006 Q4     Male   Married Full-t... 6
## 8 189891 No      36    2004      0    2007 Q1     Male   Married Full-t... 6
## 9 189891 No      36    2004      0    2007 Q2     Male   Married Full-t... 6
## 10 189894 No      33    2004      1    2010 Q3     NA     Married Full-t... 9
```

# dplyr



`dplyr` is a grammar of data manipulation, providing a consistent set of verbs that help you solve the most common data manipulation challenges:

- `mutate()` adds new variables that are functions of existing variables
- `select()` picks variables based on their names.
- `filter()` picks cases based on their values.
- `summarise()` reduces multiple values down to a single summary.
- `arrange()` changes the ordering of the rows.

These all combine naturally with `group_by()` which allows you to perform any operation “by group”.

Wickham, François, Henry, and Müller (2022)

# Transforming Variables

# mutate()

```
?dplyr::mutate
```

`mutate()` adds new variables and preserves or deletes existing ones, depending on the function arguments. New variables overwrite existing variables if they have the same name.

# mutate()

```
?dplyr::mutate
```

Function call and arguments:

```
mutate(.data,  
  ...,  
  .keep = c("all", "used",  
            "unused", "none"),  
  .before = NULL,  
  .after = NULL)
```

.data

A data.frame or tibble that should be transformed

# mutate()

```
?dplyr::mutate
```

Function call and arguments:

```
mutate(.data,  
  ...,  
  .keep = c("all", "used",  
            "unused", "none"),  
  .before = NULL,  
  .after = NULL)
```

...

- Data masking of name-value pairs
- The name specifies the variable name of the newly created column
- Multiple new variables can be created by separating each name-value pair by commas



# mutate()

```
?dplyr::mutate
```

Function call and arguments:

```
mutate(.data,  
  ...,  
  .keep = c("all", "used",  
            "unused", "none"),  
  .before = NULL,  
  .after = NULL)
```

.keep

- "all": All columns are kept in the resulting data frame (the default)
- "used": Only columns used to create new variables are kept
- "unused": Only columns that are not used to create new variables are kept
- "none": Only the newly created variables are kept

# mutate()

```
?dplyr::mutate
```

Function call and arguments:

```
mutate(.data,  
  ...,  
  .keep = c("all", "used",  
            "unused", "none"),  
  .before = NULL,  
  .after = NULL)
```

.before/.after

- Control, at which position in the data frame the new columns should be placed
- default option NULL will add the columns to the RHS

# Data Masking

- dplyr functions oftentimes use **tidy evaluation**
- tidy evaluation takes on two forms:
  - **data masking**
  - **tidy selection** (we will come to this later in this lecture)
- data masking allows for referring to variables by the name with which they are residing in an environment variable (see `vignette("programming")`)
  - Also referred to as data-variables
  - Data frame variables can be accessed simply by their name instead of the `$` operator

# Data Masking

Makes creating new variables (and other operations) more easy and intuitive:

base R

```
tbl_enoe$female <- ifelse(  
  tbl_enoe$sex == "Female",  
  "Yes",  
  "No"  
)
```

dplyr

```
mutate(  
  tbl_enoe,  
  female = ifelse(  
    sex == "Female",  
    "Yes",  
    "No"  
  )  
)
```

# Mutating a Single Variable

## Code

```
# Recode sex to a "Yes"/"No" dummy
tbl_enoe %>%
  mutate(female = ifelse(sex == "Female", "Yes", "No"),
         .before = 2)
```

## Output

```
## # A tibble: 165,457 × 13
##       id female migrate age  munic...1 fence year quarter sex  marit...2 empl_...3
##   <dbl> <chr>   <chr>   <chr> <chr>   <chr> <chr> <chr> <chr> <chr>   <chr>
## 1 189889 Yes    No      50    2004    0      2004 Q3    Fema... Single  Unempl...
## 2 189889 Yes    No      50    2004    0      2004 Q4    Fema... Single  Unempl...
## 3 189889 Yes    No      50    2004    0      2005 Q1    Fema... Single  Unempl...
## 4 189890 No     No      26    2004    0      2005 Q4    Male   Married Full-t...
## 5 189890 No     No      26    2004    0      2006 Q1    Male   Married Full-t...
## 6 189890 No     No      26    2004    0      2006 Q2    Male   <NA>    Full-t...
## 7 189891 No     No      36    2004    0      2006 Q4    Male   Married Full-t...
## 8 189891 No     No      36    2004    0      2007 Q1    Male   Married Full-t...
## 9 189891 No     No      36    2004    0      2007 Q2    Male   Married Full-t...
## 10 189894 No     No      33    2004    1      2010 Q3    NA     Married Full-t...
```

# Mutating Multiple Variables

## Code

```
# Recode marital status to a two-category variable
tbl_enoe %>%
  mutate(female = ifelse(sex == "Female", "Yes", "No"),
         married = ifelse(marital_status == "Married", "Yes", "No"),
         employed = ifelse(marital_status == "Unemployed", "No", "Yes"),
         .before = 2)
```


## Output

```
## # A tibble: 165,457 × 15
##       id female married emplo...1 migrate age  munic...2 fence year quarter sex
##   <dbl> <chr>  <chr>  <chr>  <chr>  <chr>  <chr>  <chr> <chr> <chr>  <chr>
## 1 189889 Yes    No    Yes    No    50    2004    0    2004 Q3    Fema...
## 2 189889 Yes    No    Yes    No    50    2004    0    2004 Q4    Fema...
## 3 189889 Yes    No    Yes    No    50    2004    0    2005 Q1    Fema...
## 4 189890 No     Yes   Yes    No    26    2004    0    2005 Q4    Male
## 5 189890 No     Yes   Yes    No    26    2004    0    2006 Q1    Male
## 6 189890 No     <NA> <NA>   No    26    2004    0    2006 Q2    Male
## 7 189891 No     Yes   Yes    No    36    2004    0    2006 Q4    Male
## 8 189891 No     Yes   Yes    No    36    2004    0    2007 Q1    Male
```

# Tidy Selection

Variable selection via a concise set of helper functions:

Selection helper	Description
<code>everything()</code>	Selects all columns in the data frame
<code>starts_with()/ends_with()/contains()</code>	Selects all columns starting/ending/containing a string
<code>all_of()/any_of()</code>	Selects all/any columns given by a character vector
<code>where()</code>	Selects columns based on a condition, e. g. <code>is.numeric</code>



Can be used in conjunction with `dplyr`'s `across()` function to transform multiple variables at the same time!

# Task 1: Modify the ENE/ENOE Data

Task Modify the ENE/ENOE data set for use in our analysis:

- Recode migrate, sex, marital\_status and empl\_status to meaningful 0/1 dummy variables
- Recode numerical variables that are currently stored as character columns to numerics
- Create a column containing the logarithm of income
- Recode the municipality column to a factor variable
- Create a period variable from the year and quarter columns

Code

```
tbl_enoe <- tbl_enoe %>%  
  mutate(female = ifelse(sex == "Female", "Yes", "No"), # Recode to 0/1 dummies  
         married = ifelse(marital_status == "Married", "Yes", "No"),  
         employed = ifelse(empl_status == "Unemployed", "No", "Yes"),  
         across(all_of(c("migrate", "female", "married", "employed")),  
                ~ ifelse(. == "Yes", 1, 0)),  
         # Convert all of income, educ, age and year to numerics
```



# Selecting Variables

# select()

```
?dplyr::select
```

Selects (and renames) variables in a data frame, making use of [tidy selection](#).

# select()

```
?dplyr::select
```

Function call and arguments:

```
select(.data,  
      ...)
```

`.data`

A `data.frame` or `tibble` from which columns should be selected

## Selecting Variables by Names

```
tbl_enoe %>%  
  select(id, migrate, marital_status, empl_status, educ, income, female)
```

```
## # A tibble: 165,457 × 7  
##       id migrate marital_status empl_status educ income female  
##   <dbl>   <dbl> <chr>           <chr>      <dbl> <dbl> <dbl>  
## 1 189889     0 Single      Unemployed    12     0     1  
## 2 189889     0 Single      Unemployed    12     0     1  
## 3 189889     0 Single      Unemployed    12     0     1  
## 4 189890     0 Married    Full-time     10    NA     0  
## 5 189890     0 Married    Full-time     10    NA     0  
## 6 189890     0 <NA>      Full-time     10    NA     0  
## 7 189891     0 Married    Full-time      6 3440     0  
## 8 189891     0 Married    Full-time      6 3440     0  
## 9 189891     0 Married    Full-time      6 3440     0  
## 10 189894     0 Married    Full-time      9  559     0  
## # ... with 165,447 more rows
```

# Selecting Variables by Index

```
tbl_enoe %>%  
  select(1, 2, 8:12)
```

```
## # A tibble: 165,457 × 7  
##       id migrate sex   marital_status empl_status educ income  
##   <dbl>   <dbl> <chr>   <chr>           <chr>      <dbl> <dbl>  
## 1 189889     0 Female Single      Unemployed    12     0  
## 2 189889     0 Female Single      Unemployed    12     0  
## 3 189889     0 Female Single      Unemployed    12     0  
## 4 189890     0 Male   Married      Full-time     10    NA  
## 5 189890     0 Male   Married      Full-time     10    NA  
## 6 189890     0 Male   <NA>         Full-time     10    NA  
## 7 189891     0 Male   Married      Full-time     6    3440  
## 8 189891     0 Male   Married      Full-time     6    3440  
## 9 189891     0 Male   Married      Full-time     6    3440  
## 10 189894    0 NA      Married      Full-time     9    559  
## # ... with 165,447 more rows
```

# Selecting Variables using Tidy Selection

```
tbl_enoe %>%  
  select(all_of(c("female", "employed")), ends_with("income"), where(is.factor))
```

```
## # A tibble: 165,457 × 5  
##   female employed income ln_income municipality  
##   <dbl>    <dbl>  <dbl>    <dbl>    <fct>  
## 1      1      0      0    -Inf      2004  
## 2      1      0      0    -Inf      2004  
## 3      1      0      0    -Inf      2004  
## 4      0      1     NA     NA      2004  
## 5      0      1     NA     NA      2004  
## 6      0      1     NA     NA      2004  
## 7      0      1  3440    8.14    2004  
## 8      0      1  3440    8.14    2004  
## 9      0      1  3440    8.14    2004  
## 10     0      1   559    6.33    2004  
## # ... with 165,447 more rows
```

# Selecting and Renaming Variables

```
tbl_enoe %>%  
  select(ID = id, empl = employed)
```

```
## # A tibble: 165,457 × 2  
##       ID      empl  
##   <dbl> <dbl>  
## 1 189889      0  
## 2 189889      0  
## 3 189889      0  
## 4 189890      1  
## 5 189890      1  
## 6 189890      1  
## 7 189891      1  
## 8 189891      1  
## 9 189891      1  
## 10 189894      1  
## # ... with 165,447 more rows
```

# Filtering



# `filter()`

```
?dplyr::filter
```

`filter()` subsets data frames and keeps all rows that satisfy one or more specified conditions. Applied on the rows, the condition(s) must produce the logical `TRUE` for the row to be kept in the data frame. If a condition evaluates to `NA`, the corresponding row is dropped.

# filter()

```
?dplyr::filter
```

Function call and arguments:

```
filter(.data,  
      ...)
```

.data

A `data.frame` or tibble that should be filtered

# filter()

```
?dplyr::filter
```

Function call and arguments:

```
filter(.data,  
      ...,)
```

...

- A data masking expression that returns a logical value and that is defined in terms of the variables in the `.data` argument
- Multiple expressions can be combined by the "and" (`&`) and/or "or" (`|`) operator

# Filter by a Single Variable

```
tbl_enoe %>%  
  filter(married == 1)
```

```
## # A tibble: 94,973 × 17  
##       id migrate  age munici...1 fence  year quarter sex  marit...2 empl_...3 educ  
##   <dbl>   <dbl> <dbl> <fct>   <chr> <dbl>   <dbl> <chr> <chr>   <chr>   <dbl>  
## 1 189890     0    26 2004     0    2005     3 Male Married Full-t... 10  
## 2 189890     0    26 2004     0    2006     0 Male Married Full-t... 10  
## 3 189891     0    36 2004     0    2006     3 Male Married Full-t...  6  
## 4 189891     0    36 2004     0    2007     0 Male Married Full-t...  6  
## 5 189891     0    36 2004     0    2007     1 Male Married Full-t...  6  
## 6 189894     0    33 2004     1    2010     2 NA    Married Full-t...  9  
## 7 189894     0    33 2004     1    2011     0 Fema... Married Full-t...  9  
## 8 189894     0    33 2004     1    2011     1 NA    Married Full-t...  9  
## 9 189895     0    32 2004     1    2012     0 Male Married Full-t...  9  
## 10 189895     0    32 2004     1    2012     1 Male Married Full-t...  9  
## # ... with 94,963 more rows, 6 more variables: income <dbl>, female <dbl>,  
## #   married <dbl>, employed <dbl>, ln_income <dbl>, period <dbl>, and  
## #   abbreviated variable names 1municipality, 2marital_status, 3empl_status
```

# Filter by Multiple Variables

```
tbl_enoe %>%  
  filter(married == 1, income > 0 & income <= 50000)
```

```
## # A tibble: 51,546 × 17  
##       id migrate  age munici...1 fence  year quarter sex  marit...2 empl_...3 educ  
##   <dbl>   <dbl> <dbl> <fct>   <chr> <dbl>   <dbl> <chr> <chr>   <chr>   <dbl>  
## 1 189891     0    36 2004     0    2006     3 Male Married Full-t... 6  
## 2 189891     0    36 2004     0    2007     0 Male Married Full-t... 6  
## 3 189891     0    36 2004     0    2007     1 Male Married Full-t... 6  
## 4 189894     0    33 2004     1    2010     2 NA    Married Full-t... 9  
## 5 189894     0    33 2004     1    2011     0 Fema... Married Full-t... 9  
## 6 189894     0    33 2004     1    2011     1 NA    Married Full-t... 9  
## 7 189895     0    32 2004     1    2012     0 Male Married Full-t... 9  
## 8 189895     0    32 2004     1    2012     1 Male Married Full-t... 9  
## 9 189895     0    32 2004     1    2012     2 Male Married Full-t... 9  
## 10 189902     0    30 2004     1    2010     2 Male Married Full-t... 8  
## # ... with 51,536 more rows, 6 more variables: income <dbl>, female <dbl>,  
## #   married <dbl>, employed <dbl>, ln income <dbl>, period <dbl>, and  
## #   abbreviated variable names 1municipality, 2marital_status, 3empl_status
```

# Filter using Tidy Selection

`if_any()` Filter if **any** variable satisfies the condition:

```
tbl_enoe %>%  
  filter(if_any(all_of(c("income", "employed")), ~ !is.na(.)))
```

```
## # A tibble: 165,416 × 17  
##       id migrate  age munici...1 fence  year quarter sex  marit...2 empl_...3 educ  
##   <dbl>   <dbl> <dbl> <fct>   <chr> <dbl>   <dbl> <chr> <chr>   <chr>   <dbl>  
## 1 189889     0    50 2004     0    2004     2 Fema... Single Unempl... 12  
## 2 189889     0    50 2004     0    2004     3 Fema... Single Unempl... 12  
## 3 189889     0    50 2004     0    2005     0 Fema... Single Unempl... 12  
## 4 189890     0    26 2004     0    2005     3 Male   Married Full-t... 10  
## 5 189890     0    26 2004     0    2006     0 Male   Married Full-t... 10  
## 6 189890     0    26 2004     0    2006     1 Male   <NA>    Full-t... 10  
## 7 189891     0    36 2004     0    2006     3 Male   Married Full-t...  6  
## 8 189891     0    36 2004     0    2007     0 Male   Married Full-t...  6  
## 9 189891     0    36 2004     0    2007     1 Male   Married Full-t...  6  
## 10 189894     0    33 2004     1    2010     2 NA     Married Full-t...  9
```

`if_all()` Filter if **all** variables satisfy the condition:

## Task 2: Filter the ENE/ENOE Data and Keep a Subset of Columns

### Task

1. Select all columns that are not of type "character" except `year` and `quarter`.
2. Remove all rows from the resulting data frame that contain any missing values in all columns except of `income` and `ln_income`.

### Code

```
tbl_enoe <- tbl_enoe %>%  
  select(!where(is.character), -year, -quarter) %>%  
  filter(if_all(!ends_with("income"), ~ !is.na(.)))  
  
tbl_enoe
```

### Output

```
## # A tibble: 157,248 × 11  
##       id migrate  age municipality  educ income female married emplo...1 ln_in...2  
##   <dbl>   <dbl> <dbl> <fct>          <dbl> <dbl> <dbl>   <dbl> <dbl> <dbl>
```

# **Arrange Observations**



# arrange ()

```
?dplyr::arrange
```

Orders the rows in a data frame by the values of the columns provided in the function arguments.



Useful if observations should be sorted, e. g. in a descriptive summary table!

# Arranging Observations by Income

```
tbl_enoe %>%  
  arrange(income)
```

```
## # A tibble: 157,248 × 11  
##       id migrate    age municipality educ income female married emplo...1 ln_in...2  
##   <dbl>   <dbl> <dbl> <fct>         <dbl> <dbl> <dbl>   <dbl>   <dbl>   <dbl>  
## 1 189889     0    50 2004           12     0     1     0     0    -Inf  
## 2 189889     0    50 2004           12     0     1     0     0    -Inf  
## 3 189889     0    50 2004           12     0     1     0     0    -Inf  
## 4 189898     0    28 2004            9     0     1     1     0    -Inf  
## 5 189898     0    28 2004            9     0     1     1     0    -Inf  
## 6 189898     0    28 2004            9     0     1     1     0    -Inf  
## 7 189899     0    27 2004            6     0     1     1     0    -Inf  
## 8 189908     0    15 2004            9     0     0     0     0    -Inf  
## 9 189909     0    16 2004            9     0     1     0     0    -Inf  
## 10 189909     0    16 2004            9     0     1     0     0    -Inf  
## # ... with 157,238 more rows, 1 more variable: period <dbl>, and abbreviated  
## #   variable names 1employed, 2ln_income
```

# Arranging Observations by Sex and Income in Descending Order

```
tbl_enoe %>%  
  arrange(female, desc(income))
```

```
## # A tibble: 157,248 × 11  
##       id migrate  age municipal...1 educ income female married emplo...2 ln_in...3  
##   <dbl>   <dbl> <dbl> <fct>         <dbl>   <dbl>   <dbl>   <dbl>   <dbl>   <dbl>  
## 1 5890042     0    60 28022          3 7.50e5     0     1     1    13.5  
## 2 5892262     0    41 28027          12 1.72e5     0     1     1    12.1  
## 3 5912451     0    55 28033          19 1.67e5     0     1     1    12.0  
## 4 5886803     0    33 28022          12 1.29e5     0     1     1    11.8  
## 5 5886803     0    33 28022          12 1.29e5     0     1     1    11.8  
## 6 5924420     0    42 28040          17 1.29e5     0     1     1    11.8  
## 7 376898     0    53 2002           17 1.12e5     0     1     1    11.6  
## 8 218665     0    36 2004           12 1.03e5     0     1     1    11.5  
## 9 390155     0    53 2002           16 1.00e5     0     1     1    11.5  
## 10 390155     0    53 2002           16 1.00e5     0     1     1    11.5  
## # ... with 157,238 more rows, 1 more variable: period <dbl>, and abbreviated  
## #   variable names 1municipality, 2employed, 3ln_income
```

# **Grouping and Summarizing Observations**

# group\_by()

```
?dplyr::group_by
```

- Groups data based on the specified columns and returns a grouped `tibble`
- Very useful in combination with `summary()` to create grouped summary statistics

# summarize()

```
?dplyr::summarize
```

- Creates a data frame with one (or more) rows for each combination of grouping variables
- Without grouping variables, the output has one row summarizing all observations
- Output contains one column for each specified summary statistic
- Tidy selection can be used to select columns to summarize

# Useful Summary Statistics

Function	Description
<code>mean()</code>	Returns the mean of a numerical vector.
<code>median()</code>	Returns the median of a numerical vector. See above for handling missing values.
<code>sd()</code>	Returns the standard deviation of a numerical vector.
<code>min()/max()</code>	Returns the minimum/maximum value of a numerical vector.
<code>n()</code>	Counts the numbers of observations.



Be sure to specify `na.rm = TRUE` if the vector has missing values, otherwise NA will be returned!

# Count Observations in Each Municipality

```
tbl_enoe %>%  
  group_by(municipality) %>%  
  summarise(n = n())
```

```
## # A tibble: 24 × 2  
##   municipality      n  
##   <fct>          <int>  
## 1 2004           63908  
## 2 NA             1135  
## 3 2002           35176  
## 4 2003           4731  
## 5 5025           2623  
## 6 5002           2893  
## 7 5014           649  
## 8 8037          10485  
## 9 8005            25  
## 10 26055         3021  
## # ... with 14 more rows
```



# Check for Missings in Income for Each Pair of Municipality and Period

```
tbl_enoe %>%  
  group_by(municipality, period) %>%  
  summarise(na_income = sum(is.na(income)))
```

```
## `summarise()` has grouped output by 'municipality'. You can override using the  
## `.groups` argument.
```

```
## # A tibble: 782 × 3  
## # Groups:   municipality [24]  
##   municipality period na_income  
##   <fct>         <dbl>     <int>  
## 1 2004         0.5         51  
## 2 2004         0.75        41  
## 3 2004         1          44  
## 4 2004         1.25        38  
## 5 2004         1.5         38  
## 6 2004         1.75        33  
## 7 2004         2          33
```

# Create Summary Statistics for Dummy Variables

```
tbl_enoe %>%  
  summarise(across(c("female", "married", "employed"), ~ mean(.)))
```

```
## # A tibble: 1 × 3  
##   female married employed  
##   <dbl>   <dbl>   <dbl>  
## 1  0.502   0.576   0.623
```

# Create Summary Statistics for Continuous Variables

```
tbl_enoe %>%  
  summarise(  
    across(  
      c("income", "age", "educ"),  
      .fns = list(mean = ~ mean(., na.rm = T), sd = ~ sd(., na.rm = T))  
    )  
  )
```

```
## # A tibble: 1 × 6  
##   income_mean income_sd age_mean age_sd educ_mean educ_sd  
##   <dbl>      <dbl>   <dbl>  <dbl>   <dbl>   <dbl>  
## 1      3449.      5430.    35.1   13.6     9.19    4.10
```

# Tidy Data

# What makes data tidy?

- Data can be represented in many ways:
  - Variables values may be spread over several columns, e. g. one column for each year
  - Many variables may be stored in one column, e. g. `income` and `age` are stored in the same column `value` and another column specifies which variable the value in `value` corresponds to
  - Observations may be spread across columns
- Tidy data is an organizational framework that ensure that data is stored in the correct format, i. e. it follows these rules:
  1. Variables stored in separate columns.
  2. Rows uniquely identify observations.
  3. All values are stored in their own cell.

## Example for Messy Data

```
tbl_enoe %>%  
  summarise(  
    across(  
      c("income", "age", "educ"),  
      .fns = list(  
        mean = ~ mean(., na.rm = T),  
        sd = ~ sd(., na.rm = T)  
      )  
    )  
  )  
)
```

```
## # A tibble: 1 × 6  
##   income_mean income_sd age_mean age_sd educ_mean educ_sd  
##   <dbl>      <dbl>   <dbl>  <dbl>   <dbl>   <dbl>  
## 1      3449.      5430.    35.1   13.6     9.19    4.10
```

## Example for Messy Data


```
## # A tibble: 1 × 6
##   income_mean income_sd age_mean age_sd educ_mean educ_sd
##   <dbl>      <dbl>   <dbl>  <dbl>   <dbl>   <dbl>
## 1      3449.      5430.    35.1   13.6     9.19    4.10
```

- The observation here corresponds to the column in our data set that is summarized
- The variables are the summary statistics ("mean" and "sd")



Data is stored in a too wide format!

# tidyr



The goal of `tidyr` is to help you create tidy data. Tidy data is data where:

1. Every column is variable.
2. Every row is an observation.
3. Every cell is a single value.

Tidy data describes a standard way of storing data that is used wherever possible throughout the tidyverse. If you ensure that your data is tidy, you'll spend less time fighting with the tools and more time working on your analysis.

Wickham and Girlich (2022)



# pivot\_longer

```
?tidyr::pivot_longer
```

`pivot_longer()` "lengthens" data, increasing the number of rows and decreasing the number of columns. The inverse transformation is `pivot_wider()`.

Function call and arguments:

```
pivot_longer(data,  
             cols,  
             names_to = "name",  
             names_sep = NULL,  
             values_to = "value")
```

`cols`

Tidy selection of columns to restructure into long format

# pivot\_longer

```
?tidyr::pivot_longer
```

`pivot_longer()` "lengthens" data, increasing the number of rows and decreasing the number of columns. The inverse transformation is `pivot_wider()`.

Function call and arguments:

```
pivot_longer(data,  
             cols,  
             names_to = "name",  
             names_sep = NULL,  
             values_to = "value")
```

`names_to`

- Character vector specifying the new column(s) that are created when pivoting from wide to long format
- If more than one column are created, `names_sep` (or `names_pattern`) has to be specified as well

# pivot\_longer

```
?tidyr::pivot_longer
```

`pivot_longer()` "lengthens" data, increasing the number of rows and decreasing the number of columns. The inverse transformation is `pivot_wider()`.

Function call and arguments:

```
pivot_longer(data,  
             cols,  
             names_to = "name",  
             names_sep = NULL,  
             values_to = "value")
```

`names_sep`

Either a numeric vector that specifies the position to separate the name on or a single string that specifies a regular expression to separate the name

# pivot\_longer

```
?tidyr::pivot_longer
```

`pivot_longer()` "lengthens" data, increasing the number of rows and decreasing the number of columns. The inverse transformation is `pivot_wider()`.

Function call and arguments:

```
pivot_longer(data,  
             cols,  
             names_to = "name",  
             names_sep = NULL,  
             values_to = "value")
```

`values_to`

A character vector of length 1 that specifies the column name in which to store the value

# Pivoting Summary Statistics from Wide to Long Format

## Code

```
tbl_enoe %>%  
  summarise(across(c("income", "age", "educ"),  
    .fns = list(mean = ~ mean(., na.rm = T), sd = ~ sd(., na.rm = T)))) %>%  
  pivot_longer(everything(),  
    names_to = c("variable", "statistic"),  
    names_sep = "_",  
    values_to = "value")
```

## Output

```
## # A tibble: 6 × 3  
##   variable statistic  value  
##   <chr>      <chr>    <dbl>  
## 1 income    mean      3449.  
## 2 income    sd        5430.  
## 3 age      mean       35.1  
## 4 age      sd         13.6  
## 5 educ     mean        9.19  
## 6 educ     sd          4.10
```

# pivot\_wider

```
?tidyr::pivot_wider
```

`pivot_wider()` "widens" data, increasing the number of columns and decreasing the number of rows. The inverse transformation is `pivot_longer()`.

Function call and arguments:

```
pivot_wider(data,  
            names_from = name,  
            values_from = value)
```

`names_from`

Tidy selection of columns to get the name of the output column

# pivot\_wider

```
?tidyr::pivot_wider
```

`pivot_wider()` "widens" data, increasing the number of columns and decreasing the number of rows. The inverse transformation is `pivot_longer()`.

Function call and arguments:

```
pivot_wider(data,  
            names_from = name,  
            values_from = value)
```

`values_from`

Tidy selection of columns to get the cell value from

# Pivoting from Long to Wide Format

## Code

```
tbl_enoe %>%
  summarise(across(c("income", "age", "educ"),
    .fns = list(mean = ~ mean(., na.rm = T), sd = ~ sd(., na.rm = T)))) %>%
  pivot_longer(everything(),
    names_to = c("variable", "statistic"),
    names_sep = "_",
    values_to = "value") %>%
  pivot_wider(names_from = statistic, values_from = value)
```

## Output

```
## # A tibble: 3 × 3
##   variable    mean    sd
##   <chr>    <dbl> <dbl>
## 1 income  3449.  5430.
## 2 age      35.1   13.6
## 3 educ      9.19   4.10
```



# Merging Data

# Appending Data Sets

## Row-wise

```
?dplyr::bind_rows
```

- Efficient implementation of base R's `cbind()` function that takes several data frames or a list of data frames as argument and returns a single row-binded data frame.
- Column names are used for matching the columns of the data frames and if there are columns missing in a data frame, these are filled with NA.
- Using the `.id` argument, the names of data frames lists are added in a new column in the resulting data frame.

## Column-wise

```
?dplyr::bind_cols
```

- Efficient implementation of base R's `rbind()` function that takes several data frames or a list of data frames as argument and returns a single column-binded data frame

# Binding Columns of Summary Statistics

## Code

```
# Variables to summarize
sum_stat_vars <- c("income", "age", "educ")

# Create data frame with means for each variable in sum_stat_vars
sum_stat_mean <- tbl_enoe %>%
  summarize(across(all_of(sum_stat_vars), ~ mean(., na.rm = T))) %>%
  pivot_longer(everything(),
               names_to = "variable",
               values_to = "mean")

# Create data frame with standard deviations for each variable in sum_stat_vars
sum_stat_sd <- tbl_enoe %>%
  summarize(across(all_of(sum_stat_vars), ~ sd(., na.rm = T))) %>%
  pivot_longer(everything(),
               names_to = "variable",
               values_to = "sd")

# Bind all columns of sum_stat_mean and the "sd" column of sum_stat_sd
bind_cols(sum_stat_mean, sum_stat_sd[, "sd"])
```

## Output

## Task 3: Import and Prepare the Fence Construction Data Sets

Task Take a look at the fence construction data sets residing in the directory `data/raw/fence_construction/csv/`. For each year, we have quarterly information on whether border fence construction started in a municipality or not.

Import all data sets at once, making use of the `purrr` packages. Then, bring the data into the correct format (incl. mutating columns as in [Task 1](#)) and append individual data sets to one data frame.

*Hint:* To get a list of all files residing in a directory, you can use the `list.files()` function.

Code

```
library(purrr)

# List all files in the directory
dir <- "data/raw/fence_construction/csv/"
dir_files <- list.files(dir)

tbl_fence <- paste0(dir, dir_files) %>% # Paste directory path and file names
  set_names(str_remove(dir_files, "\\\\.csv")) %>% # Set names to file names; remove suffix
  map(read_csv) %>% # Apply read_csv() over all file paths (returning a list of data frames)
```

# Joining Data Frames

- Combining a pair of data frames is achieved by joining them
- Observations in both data frames are matched by keys
- The data frames on the right have a unique identifier to match observations on (`id`)
- `dplyr` offers several ways to join both data frames by `id` to create a single data frame with both, `x` and `y`, columns




# Full Join


`full_join(x, y)`

1	x1		1	y1
2	x2		2	y2
3	x3		4	y4

# Left Join



# Right Join





# Inner Join

`inner_join(x, y)`

1	x1		1	y1
2	x2		2	y2
3	x3		4	y4

## Task 4: Join ENE/ENOE Data and Fence Construction Data

Task In order for us to be able to analyze the effect of border fence construction on the migration of Mexicans to the US, we need to combine information on individuals and information on when a fence was constructed. Fortunately, we have data for both.

Join the ENE/ENOE data frame with the fence construction data frame using appropriate keys to join observations on. Export the resulting data frame to the directory `data/processed/` as csv.

Code

```
# Ensure consistent data type of key variables
tbl_fence <- tbl_fence %>%
  mutate(municipality = forcats::as_factor(municipality))

# Inner join omits observations for which no information on fence construction is given.
# For the analysis, these observations are not needed.
tbl_out <- inner_join(tbl_enoe, tbl_fence, by = c("municipality", "period"))

write_csv(tbl_out, "data/processed/fence_migration.csv")

tbl_out
```

# References

Aden-Buie, G. (2018). *tidyexplain. Tidy Animated Verbs*. URL: <https://www.garrickadenbuie.com/project/tidyexplain/> (visited on Jan. 01, 2023).

Wickham, H., R. François, L. Henry, et al. (2022). *dplyr: A Grammar of Data Manipulation*. <https://dplyr.tidyverse.org>, <https://github.com/tidyverse/dplyr>.

Wickham, H. and M. Girlich (2022). *tidyr: Tidy Messy Data*. <https://tidyr.tidyverse.org>, <https://github.com/tidyverse/tidyr>.

Wickham, H. and G. Grolemund (2016). *R for data science. import, tidy, transform, visualize, and model data*. O'Reilly. URL: <https://r4ds.had.co.nz/>.